

World Models for Autonomous Driving: From Future Generation to Decision Making

Han Huang[✉], Dingkan Yang[✉], Lulu Guo[✉], *Member, IEEE*, Jing Cheng[✉], Yang Liu[✉], *Member, IEEE*,
Victor C.M. Leung[✉], *Life Fellow, IEEE*, Hong Chen[✉], *Fellow, IEEE*

Abstract—Autonomous driving systems must anticipate traffic evolution, compare candidate actions, and operate safely in closed loop. Recent world models address this need by generating future scenes, rolling out structured states, or exposing transitions for decision making. This survey takes a model predictive control (MPC) perspective. We treat driving world models as transition interfaces for action-conditioned rollout, cost exposure, constraint handling, uncertainty representation, and receding-horizon replanning. The literature is organized into future world generation, planning with world models, and hybrid prediction–planning systems. Across these roles, we compare visual, bird’s-eye-view, vectorized, occupancy, LiDAR, and latent states with their planner-facing interfaces. We also synthesize datasets, simulators, metrics, benchmark protocols, computational constraints, and deployment evidence. A protocol-controlled table reports comparable results without ranking incompatible settings. The review identifies open problems in long-horizon reliability, multimodal representation, generation-to-planning transfer, efficient deployment, and foundation-model-assisted safety validation. Predicted futures should ultimately be judged by the decisions they support, not by realism alone. Resources are available at the project page.

Index Terms—Autonomous driving, world models, model predictive control, planning, simulation, closed-loop evaluation.

I. INTRODUCTION

AUTONOMOUS driving requires more than accurate recognition of the current traffic scene. A driving system must also anticipate how the environment may change under different actions, interactions, and constraints. We use *driving world model* for a learned representation that can be rolled forward in time. Its future states may support generation, simulation, evaluation, or decision making. Unlike perception or conventional agent forecasting, it models scene evolution and may expose action-conditioned consequences to a planner.

This work was supported by the National Natural Science Foundation of China under Grant No. 62305113. (*Han Huang and Dingkan Yang contributed equally to this work.*)

Han Huang, Lulu Guo, Yang Liu and Hong Chen are with the College of Electronic and Information Engineering, Tongji University, Shanghai 201804, China (e-mail: guoll21@tongji.edu.cn; yangliu25@tongji.edu.cn; chenhong2019@tongji.edu.cn).

Dingkan Yang is with the College of Intelligent Robotics and Advanced Manufacturing, Fudan University, Shanghai 200433, China, and also with the Multimedia Laboratory (MMLab), the Chinese University of Hong Kong, Hong Kong SAR, China (e-mail: dkyang20@fudan.edu.cn).

Jing Cheng is with Department of Physics, School of Physics, East China Normal University, Shanghai 200241, China (e-mail: jcheng@phy.ecnu.edu.cn).

Victor C. M. Leung is with the Artificial Intelligence Research Institute, Shenzhen MSU-BIT University, Shenzhen 518172, China, the College of Computer Science and Software Engineering, Shenzhen University, Shenzhen 518060, China, and the Department of Electrical and Computer Engineering, The University of British Columbia, Vancouver, BC V6T 1Z4, Canada (email: vleung@ieee.org).

World models now span sensor synthesis, structured scene rollout, planning, and interactive simulation [1]–[3]. Definitions and evaluation protocols remain inconsistent. A focused review must distinguish what each model predicts, how a planner consumes it, and which claims rest on open-loop, non-reactive, or reactive evidence. This motivation rests on the technical scope and evidence gap documented in the cited literature, not an unreproducible publication-count trend.

Methods differ in representation and mechanism, but their role in the driving pipeline is more consequential. Some generate futures, others evaluate trajectories, and a smaller group shares prediction with planning or closed-loop simulation. We therefore organize the field by task role rather than appearance or architecture alone.

A. Comparison to Existing Surveys

Existing surveys cover world models, video and 3-D generation, connected autonomous vehicles, occupancy, foundation models, and scenario generation [1]–[12]. Motion-planning and MPC surveys instead examine constraints, receding-horizon optimization, and vehicle control [13], [14]. Neither survey line systematically analyzes the interface between learned world models and control. World-model surveys often emphasize modality or generative architecture, whereas MPC surveys usually assume an explicit predictive model. We connect them through state prediction, action-conditioned rollout, costs, constraints, uncertainty, and replanning. Table I summarizes this position.

The same predicted future can help or mislead, depending on the interface. Video favors inspection, BEV and occupancy expose geometry, and latents allow cheap rollout. Planning evidence must also account for feedback after the ego vehicle acts. Recent benchmarks therefore test actionability and repeated execution, with different degrees of agent reactivity [15]–[19].

B. Distinction, Contributions, and Roadmap

Our taxonomy separates generating a future from using it. It also distinguishes offline prediction from interactive simulation, where executed actions change the next world state.

The contributions of this survey are summarized as follows:

- We define a primary-role taxonomy—future world generation, planning with world models, and hybrid prediction–planning—with explicit boundary and secondary-capability rules.

TABLE I

POSITIONING RELATIVE TO ADJACENT SURVEY LINES. ENTRIES DESCRIBE EACH SURVEY'S PRIMARY EMPHASIS; A BLANK CIRCLE INDICATES THAT THE DIMENSION IS NOT A CENTRAL ORGANIZING AXIS, NOT THAT THE TOPIC IS ABSENT.

Survey line	Primary emphasis	AD-specific	Task-role taxonomy	MPC interface	Evidence ladder	Quantitative protocol control
General world-model surveys [4]	Concepts, agents, and cross-domain applications	○	○	○	○	○
Embodied-AI WM survey [6]	Function, time, space, and embodied-agent uses	○	○	○	○	○
Initial AD-WM survey [1]	Early concepts and application categories	●	○	○	○	○
AD-WM survey [2]	Generation, planning, interaction, and components	●	●	○	○	○
DWM ecosystem survey [3]	Modalities, applications, datasets, and ecosystem roles	●	○	○	○	○
Generation surveys [7]–[9]	Video and physical-future generation	○	○	○	○	○
CAV WM survey [10]	Cooperative and infrastructure-assisted WMs	●	○	○	○	○
Occupancy perception survey [11]	Occupancy representation, fusion, and supervision	●	○	○	○	○
FM scenario survey [12]	Foundation-model scenario generation and testing	●	○	○	●	○
Planning and MPC surveys [13], [14]	Dynamics, constraints, optimization, and control	●	○	●	○	○
Our survey	AD-WM roles, interfaces, evidence, and deployment	●	●	●	●	●

●: central organizing axis; ○: not a central organizing axis. The comparison concerns analytical organization, not overall survey quality or coverage.

- We operationalize the MPC perspective through five evidence elements: action-conditioned rollout, cost exposure, constraint handling, uncertainty representation, and re-planning. Tables II–IV apply this rubric to representative methods.
- We connect representation and mechanism choices to computation, deployment, safety-critical fidelity, and an evaluation ladder spanning open-loop, non-reactive, reactive, and higher-fidelity evidence. Table V adds a protocol-controlled quantitative synthesis.
- We derive testable research directions for long-horizon reliability, multimodal state representation, generation-to-planning transfer, efficient deployment, and foundation-model-assisted safety validation.

Fig. 1 links the definition and taxonomy to comparison, evaluation, and open challenges. Sections II–VII follow this sequence.

II. DRIVING WORLD MODELS: SCOPE AND MPC-ORIENTED INTERFACE

A. Definition

We define a world model as a learned, temporally evolving representation of the driving environment. It can be conditioned, rolled forward, and consumed by simulation or decision modules. The state may be visual, BEV, occupancy, point-cloud, vectorized, or latent. Conditions may include observations, ego actions, trajectories, maps, language, or scene metadata. Short-horizon predictors without rollout and policies without an explicit world representation fall outside our scope.

Visual plausibility alone is insufficient. Driving rollouts must preserve relevant topology, geometry, interaction, and action effects. Conversely, a useful model need not render photorealistic images; structured or latent states qualify when they support controlled rollout and decision-relevant evaluation.

We compare models through a learned transition interface that advances an explicit or latent state and exposes predicted consequences. This abstraction makes heterogeneous representations comparable without treating them as interchangeable.

B. Formal MPC-Oriented Interface

Let $\mathbf{s}_t$ be the world state, $\mathbf{a}_t$ an ego action or motion primitive, and $\mathbf{z}_t$ exogenous context. A generic stochastic transition is

$$\begin{aligned}\mathbf{s}_{t+1} &\sim \hat{p}_\theta(\cdot \mid \mathbf{s}_t, \mathbf{a}_t, \mathbf{z}_t), \\ \mathbf{o}_{t+1} &\sim \hat{p}_\theta(\cdot \mid \mathbf{s}_{t+1}, \mathbf{z}_{t+1}),\end{aligned}\tag{1}$$

where $\mathbf{o}_{t+1}$ is an optional observation-level output. Deterministic, diffusion, autoregressive, ensemble, and latent-dynamics models instantiate different forms of $\hat{p}_\theta$.

For horizon H , an MPC-style planner evaluates candidate action sequences through

$$\begin{aligned}\mathbf{a}_{0:H-1}^* &= \arg \min_{\mathbf{a}_{0:H-1}} \mathbb{E}_{\hat{p}_\theta} \left[\sum_{k=0}^{H-1} \ell(\mathbf{s}_k, \mathbf{a}_k) + \ell_T(\mathbf{s}_H) \right], \\ \text{s.t. } \mathbf{s}_{k+1} &\sim \hat{p}_\theta(\cdot \mid \mathbf{s}_k, \mathbf{a}_k, \mathbf{z}_k), \\ h(\mathbf{s}_k, \mathbf{a}_k) &\leq 0, \quad k = 0, \dots, H-1,\end{aligned}\tag{2}$$

where ℓ and ℓ_T are running and terminal costs, and h represents feasibility or safety constraints. Equation (2) is an analytical reference rather than a claim that every reviewed method




Fig. 1. Survey roadmap from the predictive transition interface to taxonomy, comparison, evaluation, and open challenges. Formal category boundaries are defined in Section III.

explicitly solves MPC. It decomposes planner-facing world modeling into five interface elements that can be inspected across heterogeneous methods.

Action-conditioned rollout asks whether a future changes under candidate ego actions. *Cost exposure* records rewards, values, risks, progress, comfort, or rule compliance that can define ℓ or ℓ_T . *Constraint handling* checks whether geometry, collisions, traffic rules, or dynamics enter generation, filtering, or optimization. *Uncertainty representation* must be decision-usable, for example by changing a safety margin or fallback rule. Finally, *replanning evidence* requires prediction and decision to repeat after execution produces a new observation. Log replay and one-shot trajectory scoring remain open loop even when they evaluate a multi-step horizon.

We record each element at one of three evidence levels. *Explicit* support means that the element is implemented in the method and evaluated through its reported interface or protocol. *Partial* support means that the representation contains relevant information, but the element is indirect, fixed, or not validated as part of the decision loop. *Not demonstrated* means that the primary paper does not provide sufficient evidence; it does not assert that the capability is impossible. This evidence rubric separates architectural potential from demonstrated planning utility and is applied in Tables II–IV.

A representative architecture follows the same flow. Current observations, map, and route context form s_t . Candidate ego actions query $\hat{p}_\theta$ to obtain multi-step state distributions. A decision module converts those rollouts into costs and constraint checks, selects an action sequence, executes its first action,

and repeats after the next observation. Future-generation models provide states; planning-oriented models add a separable decision interface. Hybrid models share prediction and action, while interactive simulators also update the environment. This architecture anchors Section III.

Fig. 2 visualizes this representative information flow and locates the reviewed method families within it.

An autoregressive model decomposes a future sequence into ordered conditional distributions,

$$p_\theta(\mathbf{x}_{1:T} | \mathbf{c}) = \prod_{\tau=1}^T p_\theta(\mathbf{x}_\tau | \mathbf{x}_{<\tau}, \mathbf{c}), \quad (3)$$

where $\mathbf{x}_{1:T}$ denotes visual, BEV, occupancy, or trajectory tokens, and $\mathbf{c}$ is the condition. Diffusion instead uses a reverse denoising process,

$$p_\theta(\mathbf{x}_0 | \mathbf{c}) = \int p(\mathbf{x}_K) \prod_{k=1}^K p_\theta(\mathbf{x}_{k-1} | \mathbf{x}_k, \mathbf{c}) d\mathbf{x}_{1:K}, \quad (4)$$

Autoregression naturally exposes sequential rollout but accumulates exposure error. Diffusion represents multimodal futures, yet repeated denoising raises query cost. Neither mechanism guarantees cost exposure, constraints, or replanning. The factorizations therefore explain query cost, uncertainty, and long-horizon behavior rather than define another taxonomy.

C. Relation to Prediction, Simulation, and Planning

World models overlap with prediction, simulation, and planning. They extend agent forecasting toward scene-level state




Fig. 2. Representative MPC-oriented information flow for a driving world model. Candidate ego-action sequences query a predictive transition; multi-step state distributions expose costs or rewards, constraints, and decision-usable uncertainty; the optimizer executes only the first selected action before replanning from the next observation. The family overlays indicate where reviewed methods enter the loop rather than asserting that every method implements all five MPC interface elements.

transitions, but they do not by themselves define the simulator or final control objective.

Motion forecasters provide multi-agent futures, map context, and uncertainty that can support world models [20]–[25]. Traffic-risk studies add decision variables that a world model may expose. Connected-vehicle trajectory prediction supports cut-in warning, potential fields quantify lane-change risk, and multi-task learning jointly predicts lane-change probability, time, and collision-risk level [26]–[28]. These components qualify only when embedded in a scene transition under intervention. We apply the same rule to end-to-end planners [29]–[31].

D. Scope, Category Boundaries, and Comparison Dimensions

The survey assigns each method by its *primary task role*. *Future world generation* produces future observations or states without requiring planning in the main objective. *Planning with world models* consumes predicted states, transitions, rewards, or values through a separable scorer, optimizer, or policy. *Hybrid prediction–planning* shares state, objective, or executable interaction between prediction and decision. A visual output alone does not make a method hybrid, and a BEV or latent output alone does not make it planning oriented.

Task role is a continuum, so a method may have secondary capabilities outside its primary category. Each method is assigned to one primary category according to its stated objective, system dependency, and strongest evaluation evidence; secondary capabilities are recorded separately. Planning-primary methods evaluate trajectory selection as a central function. Hybrid methods remove part of the separation between prediction and decision. Boundary cases are stated in the text and tables.

The assignment can be reproduced with three questions. First, can the learned transition be evaluated without a decision module? Second, does trajectory or action selection form a central reported outcome? Third, do prediction and decision share parameters, supervision, or executable feedback that cannot be removed without changing the method? A positive answer only to the first question indicates generation. The second indicates planning-primary use, whereas the third provides hybrid evidence. This procedure prevents output modality from deciding the category.

Pure perception, forecasting without world-state rollout, policies without a world representation, and language-only reasoning remain enabling components rather than primary world models.

We compare task role, representation, conditions, mechanism, output interface, and evaluation protocol. Together, these dimensions show what the state contains, how it evolves, how another module consumes it, and what evidence supports the claimed role.

III. TASK-ORIENTED TAXONOMY

Fig. 3 organizes methods by primary task role and then by representation or mechanism. The horizontal coupling arrow denotes increasing integration between the learned transition and decision modules. Generation exposes a future state, planning consumes it through a separable interface, and hybrid systems share representations, objectives, or execution. The arrow does not indicate greater visual fidelity or dimensionality.

Tables II–IV provide a representative rather than exhaustive comparison. To make the MPC mapping auditable, all three tables use a harmonized schema: subclass/state, transition or decision interface, A/C/H/U/R evidence, and evaluation evidence. We select primary works across all three task roles




Fig. 3. A task-oriented taxonomy of autonomous-driving world models based on each method’s primary task role: future world generation, planning with world models, or hybrid prediction–planning. The bottom arrow denotes increasing integration between learned transition and decision modules, not increasing visual fidelity, modality richness, or representation dimensionality.

and major state or transition families. Each selection must also support a meaningful A/C/H/U/R assessment in the primary source. Enabling components and adjacent generators remain in the prose when they provide necessary context. Each world model appears once under its primary role; secondary capabilities and boundary cases are stated in the text or table notes.

A. Future World Generation

Future world generation asks how to represent and roll forward the environment. Camera, BEV, occupancy, vector, and LiDAR states expose different errors, so fidelity must be judged against the intended downstream use.

Fig. 4 maps trade-offs in realism, geometry, sensor fidelity, and planning compatibility. A representation becomes planning oriented only when a decision module queries its rollout.

1) *Visual World Models*: Visual models generate future sensor observations that are easy to inspect but costly to convert into planning variables. GAIA-1 predicts unified video, text, and action tokens, then decodes video through diffusion [33]. DriveDreamer conditions diffusion on maps, boxes, and text; DriveDreamer-2 first induces trajectories and layouts from prompts [34], [35]. GAIA-2 adds a continuous tokenizer, multi-camera consistency, editing, and rollout [41]. STAGE and Orbis address streaming generation and longer-horizon trajectory realism [42], [43].

These models are most useful when appearance, editing, or human inspection matters. Yet a realistic video can miss the commanded lane change or interaction. It can also violate multi-camera geometry. ACT-Bench and WorldSimBench therefore test action following and action-level usefulness in addition to visual quality [18], [44].




Fig. 4. Representation-space trade-offs for future world generation. Visual, BEV, occupancy, and LiDAR world models expose different balances among perceptual realism, geometric structure, sensor fidelity, and planning compatibility.

Planner use requires a second layer of evidence. The generated scene must expose route progress, collision geometry, and the consequence of each candidate maneuver. Multi-camera consistency should also be tested after ego motion, not only at the logged viewpoint. Otherwise, visual quality can improve while the decision interface remains unchanged.

2) *BEV World Models*: BEV aligns multi-view observations on the ground plane and preserves road topology, object layout, and motion. FIERY and PowerBEV predict future instances in this space [32], [45]. StretchBEV adds stochastic residual dynamics, while CarFormer uses object-centric slots [46], [47]. MILE couples BEV latent dynamics to policy learning, and BEVWorld rolls out action-conditioned BEV tokens [48], [49].

TABLE II
REPRESENTATIVE AUTONOMOUS-DRIVING WORLD MODELS (PART 1: FUTURE WORLD GENERATION)

Study	Subclass / State	Transition and Output	MPC Evidence A/C/H/U/R	Evaluation Evidence
FIERY [32] (ICCV'21)	BEV prediction	Probabilistic BEV dynamics; future instances, motion, and trajectories.	–/–/–/E/–	nuScenes/Lyft; open-loop future prediction and multimodality.
GAIA-1 [33] (arXiv'23)	Visual / AR	Video, text, and ego-action tokens; future camera video.	E/–/P/P/–	Internal driving data; controllable visual generation, without planning evidence.
DriveDreamer [34] (ECCV'24)	Visual / diffusion	HD-map, box, and ego-action conditions; future video and layout.	E/–/P/P/–	nuScenes; generation quality, controllability, and perception augmentation.
DriveDreamer-2 [35] (AAAI'25)	Visual / diffusion	Prompt-derived trajectories and layouts; customized multi-view video.	P/–/P/P/–	nuScenes; visual quality, rare-scene control, and perception augmentation.
Copilot4D [36] (ICLR'24)	LiDAR / diffusion	Tokenized point-cloud history; future point-cloud tokens.	–/–/–/E/–	nuScenes/KITTI/AV2; open-loop Chamfer-distance forecasting.
OccWorld [37] (ECCV'24)	Occupancy / AR	Joint future occupancy and ego-trajectory tokens.	–/P/P/P/–	nuScenes; occupancy forecasting and open-loop planning metrics.
DOMe [38] (arXiv'24)	Occupancy / diffusion	Controllable multi-step occupancy generation.	P/–/P/E/–	nuScenes; long-horizon occupancy fidelity and controllability.
RangeLDM [39] (ECCV'24)	LiDAR / diffusion	Text- and layout-conditioned range-image generation.	–/–/P/E/–	nuScenes; sensor-generation fidelity, diversity, and efficiency.
SparseWorld [40] (AAAI'26)	Occupancy / sparse	Ego-state-conditioned sparse queries; 4-D occupancy futures.	P/P/P/–/–	Perception, forecasting, planning, and efficiency; no reported receding-horizon test.

MPC evidence: A, action-conditioned rollout; C, cost/reward exposure; H, constraint handling; U, decision-usable uncertainty; R, receding-horizon replanning. E denotes explicit implementation with reported interface or evaluation evidence; P denotes partial or indirect support; – denotes not demonstrated in the primary paper. The codes describe reported evidence, not theoretical capability.

BEV is compact and planner-friendly, but loses vertical structure and appearance cues. It builds on BEV perception methods that lift images, depth, temporal context, or multiple sensors into a common state [50]–[53]. World modeling adds a harder requirement: the state must remain coherent after ego and agent motion.

This coherence is inspectable. Centers, velocities, lane relations, route corridors, dynamic occupancy, and cost maps can be checked at every rollout step. Height, shape, traffic-light appearance, and subtle visual cues are less visible. BEV systems should therefore state which missing variables are recovered elsewhere in the planning stack.

3) *Vectorized and Graph-Structured World States:* Vector states encode lanes, routes, controls, and agents as polylines and relations. VectorNet established this representation for forecasting, while VAD uses vectorized agents and maps as planning constraints [20], [31]. These methods provide important state and constraint interfaces, but a vectorized planner or predictor is not automatically a world model. It must also learn how the structured state evolves under time and intervention.

VectorWorld explicitly rolls the state forward by generating ego-centric lane-agent graph tiles and coupling them to a reactive non-ego policy [54]. The planner can query objects and lanes without dense sensor rendering. However, vectorization depends on map and perception quality and can omit unexpected objects or appearance cues. Vector world models are therefore promising for efficient closed-loop rollout when explicit relations and constraints matter. Sensor-space models remain necessary when perception failure, appearance, or rendering fidelity is the evaluation target.

4) *Occupancy-Centric 4-D Scene World Models:* Occupancy and 4-D scenes represent occupied, free, and evolving

space. This structure supports collision checking and trajectory costs directly.

Broader 3-D generation offers representation lessons but should not be conflated with driving dynamics. GaussianCity uses a compact BEV-Point representation to generate unbounded 3-D cities efficiently [55]. Without temporal traffic evolution or action intervention, it is a scene generator rather than a driving world model under our definition.

DOMe combines an occupancy VAE with spatio-temporal diffusion [38]. OccLLaMA and Occ-LLM instead tokenize occupancy with language and action [56], [57]. DynamicCity generates conditional semantic 4-D occupancy, while OccWorld forecasts occupancy and ego trajectories with open-loop planning evidence [37], [58]. Drive-OccWorld connects action-controllable occupancy forecasting to planning [59]. DriveDreamer4D, UNO, and SparseWorld explore 4-D reconstruction, continuous fields, and sparse queries [40], [60], [61].

Occupancy is planner-friendly because it exposes free and uncertain space, including occluded or non-box-shaped regions. Its drawbacks are expensive supervision and rollout. Occupancy alone also omits intent, so safety depends on how the planner uses semantics and uncertainty.

False free space is the most consequential failure for planning. Average IoU can hide small occupied regions near an ego trajectory, while over-conservative occupancy can block feasible motion. Evaluation should therefore stratify errors by distance to candidate paths and report whether probabilistic occupancy changes the chosen action or safety margin.

5) *LiDAR World Models:* LiDAR models preserve metric and sensor geometry, but point clouds are sparse and viewpoint-dependent. RangeLDM uses range-view latent diffusion, and OLIDM improves object control [39], [62]. Copilot4D predicts future point clouds from BEV tokens, while

TABLE III
REPRESENTATIVE AUTONOMOUS-DRIVING WORLD MODELS (PART 2: PLANNING WITH WORLD MODELS)

Study	Subclass / State	World-Model Decision Interface	MPC Evidence A/C/H/U/R	Evaluation Evidence
MP3 [66] (CVPR'21)	Trajectory evaluation / BEV	BEV map and agent futures feed interpretable planning costs.	P/E/E/P/P	Open-loop and closed-loop urban-driving evaluation.
MBIL [48] (NeurIPS'22)	Policy learning / latent	Action-conditioned latent simulator supports imitation-policy rollout.	E/P/P/P/P	Closed-loop urban simulation; the world model is primarily used for learned imagination rather than online MPC.
DTTP [67] (ICRA'24)	Tree planning / trajectory	Candidate trajectory trees condition joint prediction and learned costs.	E/E/P/E/P	nuPlan-style open-loop planning; no reactive closed-loop evidence.
Think2Drive [68] (ECCV'24)	Policy learning / latent	Action-conditioned latent transitions, rewards, and policy imagination.	E/E/P/-/P	Closed-loop CARLA-v2 policy evaluation; world-model rollouts primarily support training.
Drive-OccWorld [59] (AAAI'25)	Trajectory evaluation / occupancy	Candidate-conditioned 4-D occupancy and occupancy-based trajectory cost.	E/E/E/P/-	nuScenes/nuScenes-Occupancy/Lyft; forecasting and open-loop planning.
Drive-WM [69] (CVPR'24)	Trajectory evaluation / visual	Candidate maneuvers condition multiview futures that a separable image-reward module compares.	E/E/P/P/-	nuScenes; open-loop planning and visual-forecasting evidence.
ReSim [70] (NeurIPS'25)	Trajectory evaluation / visual	Candidate actions condition video futures that a separable Video2Reward module scores.	E/E/P/P/-	NAVSIM policy/planning selection and video-simulation evidence; no reactive execution.
AdaptiveDriver [71] (ICRA'25)	Adaptive policy / latent	Task-adapted latent dynamics and policy under distribution shift.	E/E/E/P/E	Reactive nuPlan closed-loop evaluation with online MPC-style proposal scoring.
WoTE [72] (ICCV'25)	Trajectory evaluation / BEV	Candidate trajectories query future BEV states and trajectory scores.	E/E/P/P/E	NAVSIM open-loop and reactive Bench2Drive closed-loop evaluation.
Raw2Drive [73] (NeurIPS'25)	Policy learning / latent	Action-conditioned aligned latent rollout for policy optimization.	E/E/P/-/P	Closed-loop CARLA-v2 policy evaluation; aligned world-model rollouts primarily support training.

The A/C/H/U/R definitions and E/P/- evidence levels follow Table II. Drive-OccWorld, Drive-WM, and ReSim are assigned here by their primary evaluated decision role. Their generated occupancy or video futures remain secondary capabilities; a separable rollout-to-score interface does not by itself satisfy the hybrid criterion.

LidarDM combines a latent 4-D world with raycasting [36], [63]. LiDARcrafter adds text-conditioned scene graphs, and ViDAR uses point-cloud forecasting as camera supervision [64], [65].

Evaluation must cover object quality, temporal consistency, scan physics, and downstream utility. Sensor realism alone cannot establish better driving behavior because interaction and route feedback remain separate assumptions.

Beam pattern, missing returns, reflectance, boundaries, and occlusion all affect whether a generated scan is useful for perception. Range-view diffusion and latent raycasting encode different sampling assumptions. A credible study should disclose these assumptions and test the same downstream detector or planner on real and generated sequences.

B. Planning with World Models

Planning with world models asks how imagined futures change a decision. Outputs may be rewards, values, trajectory scores, occupancy costs, or latent transitions. Unlike logged-future generation, this role must support counterfactual ego actions through candidate rollout, latent imagination, or MPC-style optimization.

Fig. 5 distinguishes trajectory evaluation, latent policy learning, and MPC-style rollout. The subclasses depend on how the decision module consumes the model, not on state modality.

1) *Trajectory Evaluation with World Models*: Trajectory evaluators score candidate consequences. MP3 and ST-P3 combine future scene estimates with trajectory scoring [66],

[74]. WoTE maps candidates to future BEV states, and Drive-OccWorld uses action-conditioned occupancy and flow [59], [72]. Drive-WM generates multiview futures under candidate maneuvers and compares them with image-based rewards, while ReSim uses a Video2Reward module to score action-conditioned video futures [69], [70]. Their visual modality does not make them hybrid because generation and decision scoring remain separable.

This separable interface retains conventional proposals, rule checks, and comfort constraints. It is interpretable, but limited by candidate coverage and off-distribution calibration. Motion-planning studies similarly show that trajectory quality depends on spatial feasibility, temporal evolution, and safety margins rather than endpoint accuracy alone [75], [76].

Traceability is a practical advantage. An unsafe score can be linked to predicted occupancy, agent distance, rule violation, or a learned reward. However, no evaluator can recover a maneuver absent from the candidate set. Candidate diversity and uncertainty tests must therefore accompany gains in mean planning score.

2) *Latent World Models for Decision Learning*: Latent world models roll dynamics, reward, or value forward without reconstructing every observation. PlaNet and the Dreamer family established this approach, while MuZero, EfficientZero, and TD-MPC use learned states for search or trajectory optimization [77]–[83].

For driving, Think2Drive uses latent dynamics as a neural simulator; PIWM models individual interactions; Raw2Drive




Fig. 5. Planning with world models. Imagined futures are consumed through a separable decision interface for trajectory evaluation, latent policy learning, or MPC-style rollout rather than treated as standalone generated samples. These subclasses are distinguished by how the rollouts enter the decision module, not by whether the predicted state is visual, BEV, occupancy, vector, or latent.

adds privileged guidance; and AdaWM adapts selectively under task or distribution shift [68], [73], [84], [85].

Compact states make repeated rollout efficient, but may discard geometry, rules, small objects, or asymmetric interaction. Driving therefore needs structured supervision or anchors to BEV, occupancy, and trajectories. Reinforcement-learning planners reach the same conclusion: the action abstraction and learned environment must respect driving constraints [86], [87].

Hidden-state validity cannot be inferred from policy return alone. A policy may compensate for a biased latent during familiar episodes and fail after a route or weather shift. Probes for lane topology, occluded agents, traffic controls, and collision margins provide a more direct test of what survives compression.

3) *Adaptive World Models for MPC Planning:* MPC-style methods embed learned prediction in an explicit planner to improve rollout, scoring, or action selection.

MPC optimizes predicted behavior under dynamics, constraints, and costs over a receding horizon [13], [14], [88]. Motion-planning reviews and crash-mitigation planners similarly stress geometric feasibility, dynamic constraints, and traffic safety [75], [89]. Learning can refine the model, safe set, value, or terminal cost while retaining an optimization interface [83], [90]. The world model enriches MPC prediction rather than replacing MPC.

PDM-C shows that a structured proposal, internal-simulation, and scoring pipeline can outperform stronger open-loop fitting [91]. AdaptiveDriver learns interpretable IDM parameters from history and lane graphs [71]. DTPP lies near the hybrid boundary because it jointly optimizes trajectory-tree prediction and learned planning costs [67].

This approach remains interpretable and close to deployed stacks, although planner design can limit candidate diversity. It also makes uncertainty operational: probabilistic occupancy, motion modes, stochastic latents, or ensembles can alter margins and fallback behavior [92], [93]. Learned interaction models can then complement classical priors that preserve constraints and interpretable costs.

The division of responsibility should be explicit. Learned transitions can model complex interaction and adapt to local traffic, while rule-based proposals and safety envelopes retain hard feasibility checks. This arrangement also supplies a fallback when the learned model is stale, uncertain, or too slow for the current replanning deadline.


C. Hybrid Prediction–Planning World Models

Hybrid methods share prediction and action more deeply, reducing the mismatch between generated variables and planner needs. Their coupled objectives also make failures harder to attribute, so evaluation must separate representation, generation, and control effects.

TABLE IV
REPRESENTATIVE AUTONOMOUS-DRIVING WORLD MODELS (PART 3: HYBRID PREDICTION-PLANNING SYSTEMS)

Study	Subclass / State	Joint Prediction-Decision Interface	MPC Evidence A/C/H/U/R	Evaluation Evidence
GenAD [94] (ECCV'24)	Shared structured latent	Instance-centric future and ego-plan generation share scene tokens.	P/E/P/P/-	nuScenes; open-loop end-to-end planning metrics.
DrivingGPT [95] (ICCV'25)	Unified AR tokens	Observations and actions are modeled in one autoregressive sequence.	E/E/P/P/-	nuPlan/NAVSIM planning and world-modeling evidence.
World4Drive [96] (ICCV'25)	Shared physical latent	Intention-conditioned latent futures support multimodal trajectory selection.	E/E/P/E/-	nuScenes/NAVSIM; open-loop planning and non-reactive simulation metrics.
DriveArena [97] (ICCV'25)	Interactive simulator	Ego actions update a generative traffic world with simulator feedback.	E/E/P/P/E	Reactive closed-loop simulation and driving evaluation.
DriveLaW [98] (arXiv'25)	Shared visual latent	Planning and video generation share plan-conditioned latent states.	E/E/P/P/-	Planning benchmarks and video-generation evaluation.
UniDrive-WM [99] (arXiv'26)	Unified VLA tokens	Language, vision, and action tokens support understanding, planning, and generation.	E/E/P/P/-	Unified-task evaluation; limited reactive closed-loop evidence.
WorldDrive [100] (arXiv'26)	Shared vision-motion state	Trajectory-conditioned generation, inherited planning representations, and future-aware reward.	E/E/P/P/-	nuScenes, NAVSIM, and NAVSIM-v2; generation and non-reactive planning evidence.

The A/C/H/U/R definitions and E/P/- evidence levels follow Table II. Hybrid denotes a shared representation, objective, bidirectional interface, or executable interaction loop; it is not determined by image/video output alone.



At least one qualifying mechanism is required; a separable generated-future evaluator remains planning-primary

Fig. 6. Qualifying mechanisms for hybrid prediction-planning world models. Hybrid systems use at least one joint prediction-decision objective or unified observation-action interface, shared latent or structured state, or executable closed-loop interaction. A separable generated-future evaluator alone remains in the planning category.

Fig. 6 shows three qualifying mechanisms: a unified observation-action objective, a shared latent or structured state, and executable closed-loop interaction. The distinction from Fig. 5 is architectural. A planning-primary method retains a separable evaluator, optimizer, or policy interface; a hybrid method removes part of that separation.

1) *Unified Observation-Action World Models*: DrivingGPT interleaves image and action tokens and trains both through one next-token objective [95]. This is different from Drive-WM and ReSim, where candidate actions condition generated videos and a separable image-based reward or Video2Reward module evaluates the resulting futures [69], [70]. We therefore

place DrivingGPT in the hybrid category and the other two under planning.

Unified tokens can capture visual consequences omitted by hand-crafted costs. They also couple failure modes: video artifacts can alter actions, and plausible frames do not ensure feasible control.

2) *Unified Latent Driving World Models*: GenAD uses an instance-centric latent for agent evolution and ego trajectories, while World4Drive selects intention-conditioned latent futures [94], [96]. WorldDrive unifies trajectory-conditioned scene generation and planning through shared vision-motion representations and a future-aware rewarder [100]. It differs

from the earlier Generalized Predictive Model, which studies predictive pretraining and transfer [101]. DriveLaW, UniDrive-WM, and DriveWorld-VLA further connect generative latents, reasoning, and actions [98], [99], [102].

Shared latents reduce the prediction–planning boundary but raise an allocation question: which variables should be shared, and which remain task-specific? Full sharing is compact but can entangle visual detail with action choice. Partial sharing may preserve geometry while allowing separate decoders. Both require closed-loop tests.

Evaluation should isolate the source of each gain. Generation fidelity, latent consistency, action quality, and closed-loop stability need separate ablations even when trained jointly. Without them, a higher driving score cannot show whether the shared representation improved prediction or whether the planner simply learned around its errors.

3) *Closed-loop Interactive World Simulators*: Interactive simulators move world modeling to the system level. DriveArena joins a traffic manager with generative observations; DrivingSphere uses 4-D occupancy for composition and rendering [97], [103]. RealEngine integrates reconstruction, traffic composition, sensor rendering, and agent assessment, while UniSim supports neural sensor simulation and scene editing [104], [105].

Their value is executable feedback after the ego vehicle changes the scene. Credibility, however, depends on state construction, rendering, background-agent behavior, and protocol design. Waymax, ScenarioNet, Nocturne, and SimNet show how logged scenes can seed multi-agent rollout [106]–[109]. World-model simulators should report intervention robustness, agent reactivity, coverage, and reproducible failures, not only visual quality or mean driving score.

Reactive evidence requires more than looping generated frames. Background agents must respond to the ego action, maps and traffic lights must update consistently, and collisions must have repeatable semantics. Fixed seeds and scenario-level failure logs are needed so that a planning regression can be reproduced rather than absorbed into one aggregate score.

IV. CROSS-PARADIGM COMPARISON

Motivated by the receding-horizon logic of MPC, this section compares the information that each world model provides to a planner that repeatedly predicts, scores, constrains, and updates ego actions.

A. Mechanism–Task Matching

Mechanisms should be compared by task role, not ranked as generators. Diffusion suits high-dimensional futures when controllability and perceptual quality matter [38], [41], [110]. Autoregressive transformers provide token interfaces for video, text, action, occupancy, and object-centric BEV states [33], [47], [56], [111]. Latent dynamics fit rollout and value estimation because repeated imagination is cheap [68], [78], [83]. Hybrid designs connect prediction and planning, but complicate optimization and evaluation.

Control profiles differ. Diffusion captures multimodality, although sampling is costly and constraints can drift. Autoregressive models tokenize many modalities but accumulate

sequential errors. Latent dynamics support cheap imagination; their compact states are harder to verify. Graph models preserve relations and maps while omitting some appearance and dense geometry. The practical choice is the simplest representation that retains variables used by the downstream decision.

B. Cross-Representation Trade-offs

Visual states support rendering and inspection, but planning requires an additional interface. BEV offers compact geometry at the cost of vertical detail. Occupancy exposes free space directly, although dense supervision and rollout are expensive. LiDAR retains metric sensor structure but remains sparse and difficult to generate. Latent states are fastest to roll forward and hardest to inspect. In MPC terms, representations differ in how directly they expose optimizer states, constraints, and costs.

Supervision changes the trade-off. Images are abundant but weakly structured. BEV and occupancy labels require detection, alignment, fusion, or annotation, while LiDAR supplies geometry without intent. Latent states may fail silently when their objectives omit safety-critical variables. Uncertainty must be actionable: samples, cell probabilities, trajectory modes, or stochastic latents should alter speed, margins, yielding, or fallback behavior. The MPC lens asks whether uncertainty changes the feasible set, safety margin, terminal condition, or candidate cost.

C. Computational and Deployment Profile

Computational cost must be measured at the planning-system level. Diffusion repeats a denoising network, while autoregression decodes states sequentially. Dense BEV and occupancy scale with resolution and horizon. Compact latents reduce per-step cost but may still evaluate many actions. For an MPC-style consumer, latency includes sensing, encoding, candidate count N , horizon H , model updates, uncertainty samples, constraints, costs, and action selection. One rollout may meet a deadline even when $N \times H$ counterfactual transitions do not. Reports should include median and tail latency, replanning rate, candidate and sample counts, horizon, memory, hardware, precision, batching, and missed deadlines. Parameter count or FLOPs alone are insufficient.

Deployment adds stricter boundaries. Distribution shift can alter both the transition and its consuming policy. Adaptation therefore needs shift detection, regression tests, and a recoverable model version. A safety-related model also needs validity bounds, calibrated confidence, runtime monitoring, timeout behavior, and fallback. ISO 26262 covers malfunction-related safety; ISO 21448 addresses functional insufficiencies; UL 4600 frames an autonomous-product safety case [112]–[114]. These standards do not certify a model architecture. The model may propose or score futures, but a safety envelope must reject infeasible, stale, shifted, or over-budget outputs.

A deployment report should connect model latency to the control budget. It should state the sensing timestamp, deadline, worst-case candidate count, uncertainty samples, and the action taken after a miss. Average inference time is insufficient

TABLE V
COMPARABILITY-CONTROLLED QUANTITATIVE EVIDENCE FOR REPRESENTATIVE WORLD-MODEL DRIVING SYSTEMS. HIGHER IS BETTER FOR ALL REPORTED METRICS. “NON-REACTIVE” MEANS BACKGROUND ACTORS FOLLOW LOGGED FUTURES; IT IS NOT FULLY INTERACTIVE CLOSED-LOOP EVIDENCE.

Protocol group	Method	Sensor input	World-model role	Reported result	Permitted interpretation / constraint
NAVSIM test set	World4Drive [96]	C	Intention-conditioned physical latent	PDMS 85.1	Same reported test set and metric as the next row; camera-only input.
NAVSIM test set	WoTE [72]	C+L	BEV future rollout and learned trajectory evaluator	PDMS 88.3	System-level comparison under the same reported test set and metric; sensors, training, and candidate construction differ.
NAVSIM navmini	DrivingGPT [95]	front-C	Joint autoregressive observation–action tokens	PDMS 82.4	Separate navmini split; do not rank against NAVSIM test-set values.
Bench2Drive, 220 short routes	WoTE [72]	C	BEV world-model evaluator added to a TCP-based agent	DS 61.71; SR 31.36%	Reactive CARLA protocol; same test contract as the next row, with different training and policy design.
Bench2Drive, 220 short routes	Raw2Drive [73]	multi-C+IMU	Dual-stream latent models for model-based RL	DS 71.36; SR 50.24%	Supports a system-level comparison under the shared route protocol, not an isolated estimate of world-model contribution.

C: camera; front-C: front-view camera; L: LiDAR; IMU: inertial measurement unit; PDMS: Predictive Driver Model Score; DS: Driving Score; SR: Success Rate.

when rare delays suppress replanning during a conflict. The same report should distinguish failures detected by runtime monitors from those discovered only after the planner acts.

D. Utility-Oriented Comparison

Generation quality and driving utility are not equivalent. FID and FVD measure distributional realism, not geometry, control following, or planning improvement [115], [116]. Utility appears when the model changes rollout, costs, constraints, or replanning. WoTE, AdaptiveDriver, and AdaWM are less visually direct but closer to trajectory selection and adaptation [71], [72], [85]. Bench2Drive exposes closed-loop failures that open-loop metrics miss, while NAVSIM provides a non-reactive middle ground [16], [17].

Table V reports only values traceable to primary papers and groups them by split and protocol. Within-group differences in sensors, training, and planner design still preclude an architecture-only ranking. Values across groups are non-comparable.

Utility comparisons should identify which decision changes and which failure becomes detectable. They should test sensitivity to shift, horizon, and candidate quality. For MPC-style use, the model must also meet repeated-rollout deadlines and translate uncertainty into costs or constraints.

V. DATASETS, SIMULATORS, EVALUATION, AND BENCHMARKS

A. Training Datasets

Here, datasets are training resources rather than evaluation protocols. DrivingDojo supports video synthesis and conditioned generation [117]. KITTI, nuScenes, Waymo, and nuPlan provide calibrated sensors, maps, poses, and trajectories for structured modeling [15], [118]–[120]. Occ3D, OpenOccupancy, Cam4DOcc, and UniOcc add dense 3-D or 4-D supervision [121]–[124]. The key issue is coverage matching: data must supply the modality, annotations, and temporal context required by the target state.

Requirements differ by state. Video needs variation in appearance and ego motion; BEV and occupancy need calibration, pose, maps, and geometric labels. Planning needs routes and expert behavior, whereas LiDAR needs aligned scans and sensor-consistent geometry. Argoverse, Waymo Open Motion, and Lyft Level 5 remain useful for interaction and map-conditioned futures [125]–[128]. Dataset reports should describe scenario distribution, route diversity, interaction difficulty, and annotation quality, not size alone.

B. Simulation Platforms and Ecosystems

Simulation platforms determine how models interact with an environment. CARLA provides configurable sensors and closed-loop execution [129]. MetaDrive and LimSim support scalable scenarios and traffic simulation [130], [131]. Waymax initializes multi-agent scenes from logs and adds route-aware states, reactive agents, and closed-loop metrics [106]. Learned ecosystems also generate observations, turning a predictor into a simulator component.

Classical simulators offer repeatable physics and explicit rules but limited visual and behavioral diversity. Learned simulators inherit real-log statistics, yet interventions far from the data manifold remain unreliable. Hybrid systems combine explicit traffic structure with learned rendering or behavior. Each platform should disclose its state, route, intervention interface, metrics, and background-agent assumptions.

C. Evaluation Metrics

Metrics must follow the output. FID, FVD, PSNR, SSIM, and human preference mainly assess visual similarity [115], [116]. Occupancy and point clouds need geometric, temporal, and object-level measures. Conditioned models need control-following tests. Planning also requires safety, progress, comfort, rules, and closed-loop success. Static displacement error can correlate weakly with driving performance [132]. Crash mitigation and risk-aware planning therefore require safety margins under interaction, not average prediction error alone [75], [76], [93].

Metrics should be layered. Perceptual and geometric tests assess appearance and spatial validity. Controllability measures condition following. Planning and safety metrics ask whether behavior improves and severe failures decline. Visual quality cannot establish planning utility, while one closed-loop score can hide causality.

For safety-critical generation, fidelity should be decomposed into five testable levels. (1) *Conditioning fidelity* tests whether the generated ego motion follows an emergency brake or avoidance command. (2) *Kinematic and geometric fidelity* tests acceleration, jerk, yaw rate, free space, and collision geometry. (3) *Interaction fidelity* tests whether surrounding agents react causally to the intervention. (4) *Risk fidelity* uses time-to-collision, post-encroachment time, minimum distance, and collision severity rather than appearance alone. (5) *Coverage and calibration* tests long-tail category frequency and severity without equating synthetic frequency with real-world risk. A model that improves average FVD or displacement error may still miss rare cut-ins, occluded pedestrians, emergency vehicles, or construction zones. Tail-sensitive metrics, scenario strata, uncertainty intervals, and failure analysis are therefore necessary complements.

D. Benchmark Protocols

Benchmark protocols define the evaluation contract: task, split, allowed inputs, simulator assumptions, metrics, and reporting rules. nuPlan combines planning data with reactive simulation and planning metrics [15]. Bench2Drive standardizes routes, interactions, and driving skills [16]. NAVSIM is a scalable non-reactive middle ground in which background actors follow logged futures, so it does not establish reciprocal traffic response [17]. ACT-Bench directly tests whether a generated future follows a specified trajectory condition. WorldSimBench instead separates perceptual assessment from action-level usefulness [18], [44]. Shared protocols matter because private settings make heterogeneous outputs difficult to compare.

The protocol determines what progress is visible. Log replay isolates prediction under fixed futures. Non-reactive simulation measures repeated ego execution with limited interaction evidence. Reactive simulation tests reciprocal behavior but depends on background-agent realism. Scenario-generation benchmarks expose controllability and rare events, not necessarily everyday driving.

Cross-paper comparison also needs a protocol identity. Results belong to the same comparison group only when the benchmark split, permitted sensors, route set, agent-reactivity assumption, and metric implementation match. Training data and planner design may still differ, so such values remain system-level evidence. When any protocol field changes, results should be reported in separate groups rather than merged into a single ranking.

For emergency avoidance and long-tail events, a credible protocol should pair rollouts from the same initial state under distinct feasible interventions. It should verify execution of each conditioning maneuver, stratify results by conflict type and severity, and report fixed seeds for stochastic generation.

Each rollout should pass the five fidelity levels above before its downstream planner outcome is counted. This counterfactual design distinguishes an intervention-sensitive world model from one that merely produces plausible-looking hazards. Future benchmarks should state their evidence target clearly: representation quality, actionability, planning improvement, interaction robustness, safety validation, or scenario-generation capability.

E. Open-Loop and Closed-Loop Evidence

Open-loop evaluation is cheap, scalable, and reproducible, but measures agreement with logged futures under fixed distributions. Once a world model enters a planner, its predictions affect ego actions and future observations. The resulting open-loop–closed-loop gap has five mechanisms. *Observational domain shift* occurs when deployed states depart from logged inputs. *Objective mismatch* occurs when imitation or reconstruction losses omit closed-loop consequences. *Interaction mismatch* occurs when background agents remain non-reactive. *Compounding error* accumulates state and execution errors over a horizon. *Interface mismatch* occurs when model uncertainty or constraints cannot alter the planner.

These mechanisms require distinct remedies rather than one aggregate “closed-loop score.” Adaptation or data aggregation can reduce observational shift. Decision-aware and multi-step objectives can reduce objective and temporal mismatch. Reactive-agent simulation can expose interaction failures, and receding-horizon updates can limit accumulation. Calibrated uncertainty, cost, and constraint interfaces can make model limitations actionable [19].

Evidence should consequently form a ladder. Open-loop tests isolate prediction, geometry, calibration, and counterfactual controllability. Non-reactive simulation tests repeated ego execution at scale but cannot establish reciprocal interaction. Reactive closed-loop simulation tests policy–environment feedback. On-road or high-fidelity replay validation checks whether conclusions survive simulator assumptions. A claim should not be transferred upward without the corresponding experiment. Balanced evaluation combines these levels and reports where improvement disappears, treating open- and closed-loop results as complementary evidence rather than a valid/invalid dichotomy.

VI. OPEN CHALLENGES AND FUTURE TRENDS

The five challenges in Fig. 7 form a dependency chain. Representation determines which geometry, agents, semantics, and uncertainty survive long-horizon rollout. Rollout quality and planner-facing interfaces determine whether generation changes a decision. Computation limits the number of futures queried before a deadline. Foundation models broaden semantic coverage but add grounding and latency risks. Deployment and safety validation depend on all four upstream properties. Progress should therefore report state validity, horizon, decision utility, latency, and safety evidence together.




Fig. 7. MPC-oriented future challenges for autonomous-driving world models. The five directions form a dependency chain from unified multimodal state and reliable long-horizon rollout to decision utility, scalable deployment, and grounded safety evidence. Validation findings feed back to representation and rollout design, and progress should be reported as a joint profile of state validity, horizon, decision utility, latency, and safety evidence.

A. Long-Horizon Reliability

Long-horizon reliability asks whether imagined futures remain stable enough for safe decisions. More generated frames are not sufficient. Errors accumulate as temporal drift, trajectory deviation, and reduced safety margins; non-reactive agents can hide interaction failures. Protocols should vary rollout length, replanning rate, feedback, agent reactivity, calibration, and constraints. They should report when each metric becomes unreliable. One test is a horizon-conditioned curve linking rollout error to safety-margin error. Another compares one shared transition with specialized short, interaction, and route-level models under the same closed-loop protocol.

Failure thresholds should be tied to the consumer. A small pixel error may be harmless, while a one-cell free-space error near the planned path can reverse a decision. Horizon studies should therefore report both state error and the first rollout step at which the planner's selected action diverges from a trusted reference.

B. Unified Multimodal World Representation

Driving states combine appearance, 3-D geometry, maps, interactions, actions, and sometimes language. Shared BEV, occupancy, or multimodal latents seek one abstraction for prediction and planning. Feature alignment alone does not ensure geometric, causal, or control consistency. A direct test perturbs appearance, sensors, and language while holding topology and collision geometry fixed, then checks whether costs and constraints remain invariant. Selective-sharing experiments should also compare fully shared and task-specific subspaces with reconstruction and closed-loop metrics. Language and VLA models add semantics but introduce hallucination, latency, grounding, and verification risks [12], [133], [134].

Embodied multimedia extends static sensor fusion toward a perception–decision–action loop across acquisition, communication, cognition, and evaluation [135]. For driving, such signals matter only when they preserve physical semantics

and improve action. The invariances should be explicit. Appearance can vary while topology and collision geometry remain fixed; intent can vary within kinematic limits. Language describes scenarios, but physical feasibility must filter them. DriveLM, DriveVLM, and Dilu illustrate language-grounded or knowledge-driven reasoning [136]–[138].

C. Bridging Generation, Planning, and Evaluation

Realistic futures do not automatically improve decisions. Benchmarks need an evidence chain from perceptual and geometric fidelity to action conditioning, calibrated risk, and closed-loop consequence. One design fixes the planner and perturbs visual quality, geometry, condition following, and uncertainty independently. A second compares paired counterfactual maneuvers from one scene using both the model and a trusted simulator or replay oracle. Decision-aware training should preserve free space, risk, intent, rule compliance, comfort, and route progress. The evidence ladder in Section V prevents a gain at one level from supporting a claim at another.

D. Efficiency, Scalability, and Deployment

World models must support online inference and repeated candidate rollout, sometimes inside a closed-loop simulator. Dense states, diffusion, long token sequences, and language modules increase latency and memory. Compact states reduce cost but can discard safety-relevant structure. Compression should therefore be judged by deadline compliance and any loss in geometry, calibration, constraints, or safety. Candidate techniques include structured latents, distillation, quantization, token reduction, sparse activation, and event-triggered computation. The useful target is an end-to-end latency–utility frontier at fixed hardware, horizon, candidate count, and safety metrics.

Fleet and domain scaling adds another constraint. Adapters can be efficient, but uncontrolled online updates weaken reproducibility and safety assurance. Future systems should separate a frozen, regression-tested safety path from updateable

predictive components. They should monitor shift, uncertainty, deadline misses, and fallback activation, while preserving versioned evidence for every configuration. This design cannot replace model validation, but it can turn silent failure into detectable system behavior.

E. Foundation Models and Multimodal Decision Making

Foundation models add language-grounded reasoning, scenario construction, and multimodal decisions. Their roles remain distinct. Encoders identify open-vocabulary agents or rules; reasoning interfaces form candidate decisions; generators specify rare conditions; VLA policies emit actions. None alone supplies an action-conditioned transition. Semantic alignment becomes decision evidence only when grounded in geometry and when it changes a trajectory, cost, constraint, uncertainty estimate, or fallback rule [139].

DriveLM and DriveVLM illustrate semantic reasoning and hierarchical planning. DriveWorld-VLA shares world-model latents with a VLA planner, relating actions to imagined scene evolution [102], [136], [137]. Language supplies intent and open-vocabulary priors; structured states provide physical variables. The transition supplies consequences, and the controller supplies actions. Joint training does not make their evidence interchangeable. A language explanation is not proof of physical prediction or safe control.

Failure modes include hallucinated entities or rules, weak spatial grounding, prompt sensitivity, stale context, and excessive latency. Validation should pair language and geometry perturbations, test explanation–action consistency, measure calibration and deadlines, and check whether a safety monitor rejects infeasible outputs. Generative testing can expand rare-scenario coverage. Before supporting a safety claim, however, synthesized cases must pass the five fidelity levels in Section V [140]–[143].

A concrete protocol should cross language perturbations with matched geometry. It can paraphrase a valid instruction, insert a misleading prompt into the same scene, and alter geometry while holding language fixed. Action consistency, violations, confidence, and monitor interventions then separate linguistic robustness from physical grounding. A red-team benchmark should also require abstention or fallback under contradictory instructions, unseen traffic-control semantics, or missing sensor evidence. Fluent explanations alone should not receive credit.

VII. CONCLUSION

This survey studies autonomous-driving world models through an MPC lens. Future world generation rolls the state forward. Planning-oriented models convert imagined consequences into decision variables. Hybrid systems share representation, objective, or execution between prediction and action. Each role makes different trade-offs among realism, geometry, controllability, utility, and efficiency.

World models should not be evaluated only as generators. A realistic future matters when a downstream system can use it stably and verify its limits. Progress requires tighter links among generation objectives, planning interfaces, long-horizon

reliability, computation, and safety-oriented evaluation. Each claim must remain within its evidence level, from open-loop fidelity to reactive closed-loop and safety validation.

REFERENCES

- [1] Y. Guan, H. Liao, Z. Li, J. Hu, R. Yuan, Y. Li, G. Zhang, and C. Xu, "World models for autonomous driving: An initial survey," *IEEE Trans. Intell. Veh.*, pp. 1–17, 2024.
- [2] T. Feng, W. Wang, and Y. Yang, "A survey of world models for autonomous driving," 2025.
- [3] S. Tu, X. Zhou, D. Liang, X. Jiang, Y. Zhang, X. Li, and X. Bai, "The role of world models in shaping autonomous driving: A comprehensive survey," 2025.
- [4] J. Ding, Y. Zhang, Y. Shang, Y. Zhang, Z. Zong, J. Feng, Y. Yuan, H. Su, N. Li, N. Sukiennik, F. Xu, and Y. Li, "Understanding world or predicting future? a comprehensive survey of world models," *ACM Comput. Surv.*, vol. 58, no. 3, pp. 1–38, 2025.
- [5] Z. Zhu, X. Wang, W. Zhao, C. Min, B. Li, N. Deng, M. Dou, Y. Wang, B. Shi, K. Wang, C. Zhang, Y. You, Z. Zhang, D. Zhao, L. Xiao, J. Zhao, J. Lu, and G. Huang, "Is sora a world simulator? a comprehensive survey on general world models and beyond," 2025.
- [6] X. Li, X. He, L. Zhang, M. Wu, X. Li, and Y. Liu, "A comprehensive survey on world models for embodied ai," 2025.
- [7] A. Fu, Y. Zhou, T. Zhou, Y. Yang, B. Gao, Q. Li, G. Wu, and L. Shao, "Exploring the interplay between video generation and world models in autonomous driving: A survey," 2024.
- [8] B. Wen, H. Xie, Z. Chen, F. Hong, and Z. Liu, "3D scene generation: A survey," 2025.
- [9] J. Chu, Z. Liu, H. Huo, J. Qin, Y. Pan, and L. Cheng, "A survey on future physical world generation for autonomous driving," in *MMAsia Workshops*, 2025, pp. 1–8.
- [10] N. Chen and X. Liu, "Research on world models for connected automated driving: Advances, challenges, and outlook," *Appl. Sci.*, vol. 15, no. 16, p. 8986, 2025.
- [11] H. Xu, J. Chen, S. Meng, Y. Wang, and L.-P. Chau, "A survey on occupancy perception for autonomous driving: The information fusion perspective," *Inf. Fusion*, vol. 116, p. 102671, 2025.
- [12] Y. Gao, M. Piccinini, Y. Zhang, D. Wang, K. Moller, R. Brusnicki, B. Zarrouki, A. Gambi, J. F. Totz, K. Storms, S. Peters, A. Stocco, B. Alrifaa, M. Pavone, and J. Betz, "Foundation models in autonomous driving: A survey on scenario generation and scenario analysis," *IEEE Open J. Intell. Transp. Syst.*, 2026.
- [13] B. Paden, M. Čap, S. Z. Yong, D. Yershov, and E. Frazzoli, "A survey of motion planning and control techniques for self-driving urban vehicles," *IEEE Trans. Intell. Veh.*, vol. 1, no. 1, pp. 33–55, 2016.
- [14] S. Yu, M. Hirche, Y. Huang, H. Chen, and F. Allgöwer, "Model predictive control for autonomous ground vehicles: A review," *Auton. Intell. Syst.*, vol. 1, no. 4, 2021.
- [15] H. Caesar, J. Kabzan, K. S. Tan, W. K. Fong, E. M. Wolff, A. H. Lang, L. Fletcher, O. Beijbom, and S. Omari, "nuPlan: A closed-loop ML-based planning benchmark for autonomous vehicles," in *CVPR ADP3 Workshop*, 2021.
- [16] X. Jia, Z. Yang, Q. Li, Z. Zhang, and J. Yan, "Bench2Drive: Towards multi-ability benchmarking of closed-loop end-to-end autonomous driving," in *NeurIPS*, vol. 37. Curran Associates, Inc., 2024, pp. 819–844.
- [17] D. Dauner, M. Hallgarten, T. Li, X. Weng, Z. Huang, Z. Yang, H. Li, I. Gilitschenski, B. Ivanovic, M. Pavone, A. Geiger, and K. Chittita, "Navsim: Data-driven non-reactive autonomous vehicle simulation and benchmarking," in *NeurIPS*, vol. 37. Curran Associates, Inc., 2024, pp. 28 706–28 719.
- [18] Y. Qin, Z. Shi, J. Yu, X. Wang, E. Zhou, L. Li, Z. Yin, X. Liu, L. Sheng, J. Shao, L. Bai, and R. Zhang, "WorldSimBench: Towards video generation models as world simulators," in *ICML*, ser. Proceedings of Machine Learning Research, vol. 267. PMLR, 2025, pp. 50 338–50 362.
- [19] S. Z. Zhao, L. Wang, H. Ruan, Y. Bao, Y. Chen, Z. Leng, A. Ravichandran, H. He, Z. Zhou, X. Han, A. Peri, Z. Huang, P. Desai, H. Christensen, J. Ma, and B. Zhou, "BridgeSim: Unveiling the OL-CL gap in end-to-end autonomous driving," 2026.
- [20] J. Gao, C. Sun, H. Zhao, Y. Shen, D. Anguelov, C. Li, and C. Schmid, "VectorNet: Encoding HD maps and agent dynamics from vectorized representation," in *CVPR*, 2020.

- [21] J. Ngiam, V. Vasudevan, B. Caine, Z. Zhang, H.-T. L. Chiang, J. Ling, R. Roelofs, A. Bewley, C. Liu, A. Venugopal, D. J. Weiss, B. Sapp, Z. Chen, and J. Shlens, "Scene Transformer: A unified architecture for predicting future trajectories of multiple agents," in *ICLR*, 2022.
- [22] S. Shi, L. Jiang, D. Dai, and B. Schiele, "MTR++: Multi-agent motion prediction with symmetric scene modeling and guided intention querying," *IEEE Trans. Pattern Anal. Mach. Intell.*, 2023.
- [23] A. Seff, B. Cera, D. Chen, M. Ng, A. Zhou, N. Nayakanti, K. S. Refaat, R. Al-Rfou, and B. Sapp, "MotionLM: Multi-agent motion forecasting as language modeling," in *ICCV*, 2023.
- [24] N. Nayakanti, R. Al-Rfou, A. Zhou, K. Goel, K. S. Refaat, and B. Sapp, "Wayformer: Motion forecasting via simple & efficient attention networks," in *ICRA*, 2023.
- [25] R. Mahjourian, J. Kim, Y. Chai, M. Tan, B. Sapp, and D. Anguelov, "Occupancy flow fields for motion forecasting in autonomous driving," *IEEE Robot. Autom. Lett.*, 2022.
- [26] N. Lyu, J. Wen, Z. Duan, and C. Wu, "Vehicle trajectory prediction and cut-in collision warning model in a connected vehicle environment," *IEEE Trans. Intell. Transp. Syst.*, vol. 23, no. 2, pp. 966–981, 2022.
- [27] Y. Wang, N. Lyu, and J. Wen, "Modeling risk potential fields for mandatory lane changes in intelligent connected vehicle environment," *Expert Syst. Appl.*, vol. 255, p. 124814, 2024.
- [28] L. Yang, J. Zhang, N. Lyu, and Q. Zhao, "Predicting lane change maneuver and associated collision risks based on multi-task learning," *Accid. Anal. Prev.*, vol. 209, p. 107830, 2025.
- [29] F. Codevilla, M. Muller, A. Lopez, V. Koltun, and A. Dosovitskiy, "End-to-end driving via conditional imitation learning," in *ICRA*, 2018.
- [30] K. Chitta, A. Prakash, B. Jaeger, Z. Yu, K. Renz, and A. Geiger, "TransFuser: Imitation with transformer-based sensor fusion for autonomous driving," *IEEE Trans. Pattern Anal. Mach. Intell.*, 2023.
- [31] B. Jiang, S. Chen, Q. Xu, B. Liao, J. Chen, H. Zhou, Q. Zhang, W. Liu, C. Huang, and X. Wang, "VAD: Vectorized scene representation for efficient autonomous driving," in *ICCV*, 2023.
- [32] A. Hu, Z. Murez, N. Mohan, S. Dudas, J. Hawke, V. Badrinarayanan, R. Cipolla, and A. Kendall, "FIERY: Future instance prediction in bird's-eye view from surround monocular cameras," in *ICCV*, 2021.
- [33] A. Hu, L. Russell, H. Yeo, Z. Murez, G. Fedoseev, A. Kendall, J. Shotton, and G. Corrado, "GAIA-1: A generative world model for autonomous driving," 2023.
- [34] X. Wang, Z. Zhu, G. Huang, X. Chen, J. Zhu, and J. Lu, "DriveDreamer: Towards real-world-drive world models for autonomous driving," in *ECCV*. Springer Nature Switzerland, 2024, pp. 55–72.
- [35] G. Zhao, X. Wang, Z. Zhu, X. Chen, G. Huang, X. Bao, and X. Wang, "DriveDreamer-2: LLM-enhanced world models for diverse driving video generation," in *AAAI*, 2025.
- [36] L. Zhang, Y. Xiong, Z. Yang, S. Casas, R. Hu, and R. Urtasun, "Copilot4D: Learning unsupervised world models for autonomous driving via discrete diffusion," in *ICLR*, 2024.
- [37] W. Zheng, W. Chen, Y. Huang, B. Zhang, Y. Duan, and J. Lu, "OccWorld: Learning a 3D occupancy world model for autonomous driving," in *ECCV*. Springer Nature Switzerland, 2024, pp. 55–72.
- [38] S. Gu, W. Yin, B. Jin, X. Guo, J. Wang, H. Li, Q. Zhang, and X. Long, "DOME: Taming diffusion model into high-fidelity controllable occupancy world model," 2024.
- [39] Q. Hu, Z. Zhang, and W. Hu, "RangeLDM: Fast realistic LiDAR point cloud generation," in *ECCV*, ser. Lecture Notes in Computer Science, 2024.
- [40] C. Dang, H. Liu, J. Bao, P. An, X. Tang, A. Pan, J. Ma, B. Sun, and Y. Wang, "SparseWorld: A flexible, adaptive, and efficient 4D occupancy world model powered by sparse and dynamic queries," in *AAAI*, vol. 40, no. 5, 2026, pp. 3497–3505.
- [41] L. Russell, A. Hu, L. Bertoni, G. Fedoseev, J. Shotton, E. Arani, and G. Corrado, "GAIA-2: A controllable multi-view generative world model for autonomous driving," 2025.
- [42] J. Wang, Y. Yao, X. Feng, H. Wu, Y. Wang, Q. Huang, Y. Ma, and X. Zhu, "STAGE: A stream-centric generative world model for long-horizon driving-scene simulation," in *IROS*, 2025.
- [43] A. Mousakhan, S. Mittal, S. Galetto, K. Farid, and T. Brox, "Orbis: Overcoming challenges of long-horizon prediction in driving world models," in *NeurIPS*, 2025.
- [44] H. Arai, K. Ishihara, T. Takahashi, and Y. Yamaguchi, "ACT-Bench: Towards action controllable world models for autonomous driving," in *ICLR 2025 Workshop World Models*, 2025. [Online]. Available: <https://arxiv.org/abs/2412.05337>
- [45] P. Li, S. Ding, X. Chen, N. Hanselmann, M. Cordts, and J. Gall, "PowerBEV: A powerful yet lightweight framework for instance prediction in bird's-eye view," in *IJCAI*, 2023, pp. 1080–1088.
- [46] A. K. Akan and F. Güney, "StretchBEV: Stretching future instance prediction spatially and temporally," in *ECCV*, ser. Lecture Notes in Computer Science, 2022.
- [47] S. Hamdan and F. Güney, "CarFormer: Self-driving with learned object-centric representations," in *ECCV*. Cham: Springer Nature Switzerland, 2024, pp. 177–193.
- [48] A. Hu, G. Corrado, N. Griffiths, Z. Murez, C. Gurau, H. Yeo, A. Kendall, R. Cipolla, and J. Shotton, "Model-based imitation learning for urban driving," in *NeurIPS*, vol. 35. Curran Associates, Inc., 2022, pp. 20 703–20 716.
- [49] Y. Zhang, S. Gong, K. Xiong, X. Ye, X. Li, X. Tan, F. Wang, J. Huang, H. Wu, and H. Wang, "Bevworld: A multimodal world simulator for autonomous driving via scene-level bev latents," 2025.
- [50] J. Philion and S. Fidler, "Lift, splat, shoot: Encoding images from arbitrary camera rigs by implicitly unprojecting to 3D," in *ECCV*. Cham: Springer International Publishing, 2020, pp. 194–210.
- [51] Z. Li, W. Wang, H. Li, E. Xie, C. Sima, T. Lu, Q. Yu, and J. Dai, "Bevformer: Learning bird's-eye-view representation from lidar-camera via spatiotemporal transformers," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 47, no. 3, pp. 2020–2036, 2025.
- [52] Y. Li, Z. Ge, G. Yu, J. Yang, Z. Wang, Y. Shi, J. Sun, and Z. Li, "BEVDepth: Acquisition of reliable depth for multi-view 3D object detection," in *AAAI*, 2023.
- [53] Z. Liu, H. Tang, A. Amini, X. Yang, H. Mao, D. L. Rus, and S. Han, "BEVFusion: Multi-task multi-sensor fusion with unified bird's-eye view representation," in *ICRA*, 2023.
- [54] C. Jiang, D. Zhou, J. Liu, and K. L. Sun, "VectorWorld: Efficient streaming world model via diffusion flow on vector graphs," 2026. [Online]. Available: <https://arxiv.org/abs/2603.17652>
- [55] H. Xie, Z. Chen, F. Hong, and Z. Liu, "Generative gaussian splatting for unbounded 3D city generation," in *CVPR*, 2025, pp. 6111–6120.
- [56] J. Wei, S. Yuan, P. Li, Q. Hu, Z. Gan, and W. Ding, "Occllama: An occupancy-language-action generative world model for autonomous driving," 2024.
- [57] T. Xu, H. Lu, X. Yan, Y. Cai, B. Liu, and Y. Chen, "Occ-LLM: Enhancing autonomous driving with occupancy-based large language models," in *ICRA*, 2025, pp. 8434–8441.
- [58] H. Bian, L. Kong, H. Xie, L. Pan, Y. Qiao, and Z. Liu, "DynamicCity: Large-scale 4D occupancy generation from dynamic scenes," in *ICLR*, 2025.
- [59] Y. Yang, J. Mei, Y. Ma, S. Du, W. Chen, Y. Qian, Y. Feng, and Y. Liu, "Driving in the occupancy world: Vision-centric 4D occupancy forecasting and planning via world models for autonomous driving," in *AAAI*, 2025.
- [60] G. Zhao, C. Ni, X. Wang, Z. Zhu, X. Zhang, Y. Wang, G. Huang, X. Chen, B. Wang, Y. Zhang, W. Mei, and X. Wang, "DriveDreamer4D: World models are effective data machines for 4D driving scene representation," in *CVPR*, 2025, pp. 12 015–12 026.
- [61] B. Agro, Q. Sykora, S. Casas, T. Gilles, and R. Urtasun, "UnO: Unsupervised occupancy fields for perception and forecasting," in *CVPR*, 2024.
- [62] T. Yan, J. Yin, X. Lang, R. Yang, C.-Z. Xu, and J. Shen, "OLiDM: Object-aware LiDAR diffusion models for autonomous driving," in *AAAI*, 2025.
- [63] V. Zyrianov, H. Che, Z. Liu, and S. Wang, "LidarDM: Generative LiDAR simulation in a generated world," in *ICRA*, 2025.
- [64] A. Liang, Y. Liu, Y. Yang, D. Lu, L. Li, L. Kong, H. Zhao, and W. T. Ooi, "LiDARcrafter: Dynamic 4D world modeling from LiDAR sequences," in *AAAI*, vol. 40, no. 22, 2026, pp. 18 406–18 414.
- [65] Z. Yang, L. Chen, Y. Sun, and H. Li, "Visual point cloud forecasting enables scalable autonomous driving," in *CVPR*, 2024.
- [66] S. Casas, A. Sadat, and R. Urtasun, "MP3: A unified model to map, perceive, predict and plan," in *CVPR*, 2021.
- [67] Z. Huang, P. Karkus, B. Ivanovic, Y. Chen, M. Pavone, and C. Lv, "DTPP: Differentiable joint conditional prediction and cost evaluation for tree policy planning in autonomous driving," in *ICRA*, 2024.
- [68] Q. Li, X. Jia, S. Wang, and J. Yan, "Think2Drive: Efficient reinforcement learning by thinking with latent world model for autonomous driving (in CARLA-v2)," in *ECCV*, ser. Lecture Notes in Computer Science. Springer, 2024, pp. 142–158.
- [69] Y. Wang, J. He, L. Fan, H. Li, Y. Chen, and Z. Zhang, "Driving into the future: Multiview visual forecasting and planning with world model for autonomous driving," in *CVPR*, 2024.
- [70] J. Yang, K. Chitta, S. Gao, L. Chen, Y. Shao, X. Jia, H. Li, A. Geiger, X. Yue, and L. Chen, "ReSim: Reliable world simulation for autonomous driving," in *NeurIPS*, 2025.

- [71] A. B. Vasudevan, N. Peri, J. Schneider, and D. Ramanan, "Planning with adaptive world models for autonomous driving," in *ICRA*, 2025.
- [72] Y. Li, Y. Wang, Y. Liu, J. He, L. Fan, and Z. Zhang, "End-to-end driving with online trajectory evaluation via BEV world model," in *ICCV*, 2025, pp. 27 137–27 146.
- [73] Z. Yang, X. Jia, Q. Li, X. Yang, M. Yao, and J. Yan, "Raw2Drive: Reinforcement learning with aligned world models for end-to-end autonomous driving (in CARLA v2)," in *NeurIPS*, 2025.
- [74] S. Hu, L. Chen, P. Wu, H. Li, J. Yan, and D. Tao, "ST-P3: End-to-end vision-based autonomous driving via spatial-temporal feature learning," in *ECCV*, ser. Lecture Notes in Computer Science, 2022.
- [75] H. Wang, Y. Huang, A. Khajepour, Y. Zhang, Y. Rasekhipour, and D. Cao, "Crash mitigation in motion planning for autonomous vehicles," *IEEE Trans. Intell. Transp. Syst.*, vol. 20, no. 9, pp. 3313–3323, 2019.
- [76] Z. Han, Y. Wu, T. Li, L. Zhang, L. Pei, L. Xu, C. Li, C. Ma, C. Xu, S. Shen, and F. Gao, "An efficient spatial-temporal trajectory planner for autonomous vehicles in unstructured environments," *IEEE Trans. Intell. Transp. Syst.*, vol. 25, no. 2, pp. 1797–1814, 2024.
- [77] D. Hafner, T. Lillicrap, I. Fischer, R. Villegas, D. Ha, H. Lee, and J. Davidson, "Learning latent dynamics for planning from pixels," in *ICML*, ser. Proceedings of Machine Learning Research, vol. 97. PMLR, 2019, pp. 2555–2565.
- [78] D. Hafner, T. Lillicrap, J. Ba, and M. Norouzi, "Dream to control: Learning behaviors by latent imagination," in *ICLR*, 2020.
- [79] D. Hafner, T. P. Lillicrap, M. Norouzi, and J. Ba, "Mastering atari with discrete world models," in *ICLR*, 2021.
- [80] D. Hafner, J. Pasukonis, J. Ba, and T. Lillicrap, "Mastering diverse control tasks through world models," *Nature*, vol. 640, pp. 647–653, 2025.
- [81] J. Schrittwieser, I. Antonoglou, T. Hubert, K. Simonyan, L. Sifre, S. Schmitt, A. Guez, E. Lockhart, D. Hassabis, T. Graepel, T. Lillicrap, and D. Silver, "Mastering atari, go, chess and shogi by planning with a learned model," *Nature*, 2019.
- [82] W. Ye, S. Liu, T. Kurutach, P. Abbeel, and Y. Gao, "Mastering atari games with limited data," in *NeurIPS*, vol. 34. Curran Associates, Inc., 2021.
- [83] N. A. Hansen, H. Su, and X. Wang, "Temporal difference learning for model predictive control," in *ICML*, ser. Proceedings of Machine Learning Research, vol. 162. PMLR, 2022, pp. 8387–8406.
- [84] Y. Gao, Q. Zhang, D. wei Ding, and D. Zhao, "Dream to drive with predictive individual world model," *IEEE Trans. Intell. Veh.*, 2024.
- [85] H. Wang, X. Ye, F. Tao, C. Pan, A. Mallik, B. Yaman, L. Ren, and J. Zhang, "AdaWM: Adaptive world model based planning for autonomous driving," in *ICLR*, 2025.
- [86] S. Aradi, "Survey of deep reinforcement learning for motion planning of autonomous vehicles," *IEEE Trans. Intell. Transp. Syst.*, vol. 23, no. 2, pp. 740–759, 2022.
- [87] T. Schneider, M. V. A. Pedrosa, T. P. Gros, V. Wolf, and K. Flaßkamp, "Motion primitives as the action space of deep q-learning for planning in autonomous driving," *IEEE Trans. Intell. Transp. Syst.*, vol. 25, no. 11, pp. 17 852–17 864, 2024.
- [88] P. Falcone, F. Borrelli, J. Asgari, H. E. Tseng, and D. Hrovat, "Predictive active steering control for autonomous vehicle systems," *IEEE Trans. Control Syst. Technol.*, vol. 15, no. 3, pp. 566–580, 2007.
- [89] D. González, J. Pérez, V. Milanés, and F. Nashashibi, "A review of motion planning techniques for automated vehicles," *IEEE Trans. Intell. Transp. Syst.*, vol. 17, no. 4, pp. 1135–1145, 2016.
- [90] U. Rosolia and F. Borrelli, "Learning how to autonomously race a car: A predictive control approach," *IEEE Trans. Control Syst. Technol.*, vol. 28, no. 6, pp. 2713–2719, 2020.
- [91] D. Dauner, M. Hallgarten, A. Geiger, and K. Chitta, "Parting with misconceptions about learning-based vehicle motion planning," in *CoRL*. PMLR, 2023, pp. 1268–1281.
- [92] J. Li, W. Zhan, Y. Hu, and M. Tomizuka, "Generic tracking and probabilistic prediction framework and its application in autonomous driving," *IEEE Trans. Intell. Transp. Syst.*, vol. 21, no. 9, pp. 3634–3649, 2020.
- [93] Q. Wang, D. Xu, G. Kuang, C. Lv, S. E. Li, and B. Nie, "Risk-aware vehicle trajectory prediction under safety-critical scenarios," *IEEE Trans. Intell. Transp. Syst.*, vol. 26, no. 5, pp. 7156–7171, 2025.
- [94] W. Zheng, R. Song, X. Guo, C. Zhang, and L. Chen, "GenAD: Generative end-to-end autonomous driving," in *ECCV*, ser. Lecture Notes in Computer Science, 2024.
- [95] Y. Chen, Y. Wang, and Z. Zhang, "DrivingGPT: Unifying driving world modeling and planning with multi-modal autoregressive transformers," in *ICCV*, 2025, pp. 26 890–26 900.
- [96] Y. Zheng, P. Yang, Z. Xing, Q. Zhang, Y. Zheng, Y. Gao, P. Li, T. Zhang, Z. Xia, P. Jia, X. Lang, and D. Zhao, "World4Drive: End-to-end autonomous driving via intention-aware physical latent world model," in *ICCV*, 2025.
- [97] X. Yang, L. Wen, T. Wei, Y. Ma, J. Mei, X. Li, W. Lei, D. Fu, P. Cai, M. Dou, L. He, Y. Liu, B. Shi, and Y. Qiao, "DriveArena: A closed-loop generative simulation platform for autonomous driving," in *ICCV*, 2025, pp. 26 933–26 943.
- [98] T. Xia, Y. Li, L. Zhou, J. Yao, K. Xiong, H. Sun, B. Wang, K. Ma, G. Chen, H. Ye, W. Liu, and X. Wang, "DriveLaW: Unifying planning and video generation in a latent driving world," 2025.
- [99] Z. Xiong, X. Ye, B. Yaman, S. Cheng, Y. Lu, J. Luo, N. Jacobs, and L. Ren, "UniDrive-WM: Unified understanding, planning and generation world model for autonomous driving," 2026.
- [100] X. Gui, M. Zhang, T. Yan, W. Han, J. Gong, F. Tan, C.-Z. Xu, and J. Shen, "Bridging scene generation and planning: Driving with world model via unifying vision and motion representation," 2026. [Online]. Available: <https://arxiv.org/abs/2603.14948>
- [101] J. Yang, S. Gao, Y. Qiu, L. Chen, T. Li, B. Dai, K. Chitta, P. Wu, J. Zeng, P. Luo, J. Zhang, A. Geiger, Y. Qiao, and H. Li, "Generalized predictive model for autonomous driving," in *CVPR*, 2024, pp. 14 662–14 672.
- [102] F. Jia, L. Liu, Z. Song, C. Jia, H. Ye, X. Hao, and L. Chen, "DriveWorld-VLA: Unified latent-space world modeling with vision-language-action for autonomous driving," 2026.
- [103] T. Yan, D. Wu, W. Han, J. Jiang, X. Zhou, K. Zhan, C.-Z. Xu, and J. Shen, "DrivingSphere: Building a high-fidelity 4D world for closed-loop simulation," in *CVPR*, 2025.
- [104] J. Jiang, N. Song, J. Li, X. Zhu, and L. Zhang, "RealEngine: Simulating autonomous driving in realistic context," 2025.
- [105] Z. Yang, Y. Chen, J. Wang, S. Manivasagam, W.-C. Ma, A. J. Yang, and R. Urtasun, "UniSim: A neural closed-loop sensor simulator," in *CVPR*, 2023.
- [106] C. Gulino, J. Fu, W. Luo, G. Tucker, E. Bronstein, Y. Lu, J. Harb, X. Pan, Y. Wang, X. Chen, J. Co-Reyes, R. Agarwal, R. Roelofs, Y. Lu, N. Montali, P. Mougins, Z. Yang, B. White, A. Faust, R. McAllister, D. Anguelov, and B. Sapp, "Waymax: An accelerated, data-driven simulator for large-scale autonomous driving research," in *NeurIPS*, vol. 36. Curran Associates, Inc., 2023, pp. 7730–7742.
- [107] C. Duan, L. Feng, Q. Li, Z. Liu, W. Mo, Z. M. Peng, and B. Zhou, "ScenarioNet: Open-source platform for large-scale traffic scenario simulation and modeling," in *NeurIPS*, 2023.
- [108] B. Amos, J. Foerster, N. Lichtle, E. Vinitsky, and X. Yang, "Nocturne: A scalable driving benchmark for bringing multi-agent learning one step closer to the real world," in *NeurIPS*, 2022.
- [109] L. Bergamini, Y. Ye, O. Scheel, L. Chen, C. Hu, L. D. Pero, B. Osinski, H. Grimmett, and P. Ondruska, "SimNet: Learning reactive self-driving simulations from real-world observations," in *ICRA*, 2021.
- [110] J. Ho, A. Jain, and P. Abbeel, "Denoising diffusion probabilistic models," in *NeurIPS*, vol. 33. Curran Associates, Inc., 2020, pp. 6840–6851.
- [111] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, "Attention is all you need," in *NIPS*, vol. 30. Curran Associates, Inc., 2017.
- [112] *Road Vehicles—Functional Safety*, International Organization for Standardization Std. ISO 26 262:2018, 2018. [Online]. Available: <https://www.iso.org/standard/68383.html>
- [113] *Road Vehicles—Safety of the Intended Functionality*, International Organization for Standardization Std. ISO 21 448:2022, 2022. [Online]. Available: <https://www.iso.org/standard/77490.html>
- [114] *Standard for Safety for the Evaluation of Autonomous Products*, UL Standards and Engagement Std. ANSI/UL 4600, 2023. [Online]. Available: <https://www.ul.com/news/ul-4600-edition-3-updates-incorporate-autonomous-trucking>
- [115] M. Heusel, H. Ramsauer, T. Unterthiner, B. Nessler, and S. Hochreiter, "Gans trained by a two time-scale update rule converge to a local nash equilibrium," in *NIPS*, vol. 30. Curran Associates, Inc., 2017.
- [116] T. Unterthiner, S. van Steenkiste, K. Kurach, R. Marinier, M. Michalski, and S. Gelly, "Towards accurate generative models of video: A new metric & challenges," 2019.
- [117] Y. Wang, K. Cheng, J. He, Q. Wang, H. Dai, Y. Chen, F. Xia, and Z. Zhang, "Drivingdojo dataset: Advancing interactive and knowledge-enriched driving world model," in *NeurIPS*, vol. 37. Curran Associates, Inc., 2024, pp. 13 020–13 034.
- [118] A. Geiger, P. Lenz, and R. Urtasun, "Are we ready for autonomous driving the KITTI vision benchmark suite," in *CVPR*, 2012.

- [119] H. Caesar, V. Bankiti, A. H. Lang, S. Vora, V. E. Liong, Q. Xu, A. Krishnan, Y. Pan, G. Baldan, and O. Beijbom, “nuScenes: A multimodal dataset for autonomous driving,” in *CVPR*, 2020.
- [120] P. Sun, H. Kretschmar, X. Dotiwalla, A. Chouard, V. Patnaik, P. Tsui, J. Guo, Y. Zhou, Y. Chai, B. Caine, V. Vasudevan, W. Han, J. Ngiam, H. Zhao, A. Timofeev, S. Ettinger, M. Krivokon, A. Gao, A. Joshi, Y. Zhang, J. Shlens, Z. Chen, and D. Anguelov, “Scalability in perception for autonomous driving: Waymo open dataset,” in *CVPR*, 2020.
- [121] X. Tian, T. Jiang, L. Yun, Y. Mao, H. Yang, Y. Wang, Y. Wang, and H. Zhao, “Occ3d: A large-scale 3d occupancy prediction benchmark for autonomous driving,” in *NeurIPS*, vol. 36. Curran Associates, Inc., 2023, pp. 64 318–64 330.
- [122] X. Wang, Z. Zhu, W. Xu, Y. Zhang, Y. Wei, X. Chi, Y. Ye, D. Du, J. Lu, and X. Wang, “OpenOccupancy: A large scale benchmark for surrounding semantic occupancy perception,” in *ICCV*, 2023.
- [123] J. Ma, X. Chen, J. Huang, J. Xu, Z. Luo, J. Xu, W. Gu, R. Ai, and H. Wang, “Cam4DOcc: Benchmark for camera-only 4D occupancy forecasting in autonomous driving applications,” in *CVPR*, 2024.
- [124] Y. Wang, X. Huang, X. Sun, M. Yan, S. Xing, Z. Tu, and J. Li, “UniOcc: A unified benchmark for occupancy forecasting and prediction in autonomous driving,” in *ICCV*, 2025, pp. 25 560–25 570.
- [125] M.-F. Chang, J. Lambert, P. Sangkloy, J. Singh, S. Bak, A. Hartnett, D. Wang, P. Carr, S. Lucey, D. Ramanan, and J. Hays, “Argoverse: 3D tracking and forecasting with rich maps,” in *CVPR*, 2019.
- [126] B. Wilson, W. Qi, T. Agarwal, J. Lambert, J. Singh, S. Khandelwal, B. Pan, R. Kumar, A. Hartnett, J. K. Pontes, D. Ramanan, P. Carr, and J. Hays, “Argoverse 2: Next generation datasets for self-driving perception and forecasting,” in *NeurIPS Datasets and Benchmarks Track*, 2021.
- [127] S. Ettinger, S. Cheng, B. Caine, C. Liu, H. Zhao, S. Pradhan, Y. Chai, B. Sapp, C. R. Qi, Y. Zhou, Z. Yang, A. Chouard, P. Sun, J. Ngiam, V. Vasudevan, A. McCauley, J. Shlens, and D. Anguelov, “Large scale interactive motion forecasting for autonomous driving: The waymo open motion dataset,” in *ICCV*, 2021.
- [128] G. Li, Y. Jiao, V. L. Knoop, S. C. Calvert, and J. Van Lint, “Large car-following data based on lyft level-5 open dataset: Following autonomous vehicles vs. human-driven vehicles,” in *ITSC*, 2023, pp. 5818–5823.
- [129] A. Dosovitskiy, G. Ros, F. Codevilla, A. Lopez, and V. Koltun, “CARLA: An open urban driving simulator,” in *CoRL*, ser. Proceedings of Machine Learning Research, vol. 78. PMLR, 2017, pp. 1–16.
- [130] Q. Li, Z. Peng, L. Feng, Q. Zhang, Z. Xue, and B. Zhou, “MetaDrive: Composing diverse driving scenarios for generalizable reinforcement learning,” *IEEE Trans. Pattern Anal. Mach. Intell.*, 2022.
- [131] L. Wen, D. Fu, S. Mao, P. Cai, M. Dou, Y. Li, and Y. Qiao, “LimSim: A long-term interactive multi-scenario traffic simulator,” in *ITSC*, 2023.
- [132] T. Phong, H. Wu, C. Yu, P. Cai, S. Zheng, and D. Hsu, “What truly matters in trajectory prediction for autonomous driving?” in *NeurIPS*, vol. 36. Curran Associates, Inc., 2023, pp. 71 327–71 339.
- [133] S. Jiang, Z. Huang, K. Qian, Z. Luo, T. Zhu, Y. Zhong, Y. Tang, M. Kong, Y. Wang, S. Jiao, H. Ye, Z. Sheng, X. Zhao, T. Wen, Z. Fu, S. Chen, K. Jiang, D. Yang, S. Choi, and L. Sun, “A survey on vision-language-action models for autonomous driving,” in *ICCVW*, 2025.
- [134] K. Yu, Y. Sun, T. Wu, R. Zhang, Z. Lin, Y. Meng, J. Wang, and Y. Yang, “A survey of reasoning in autonomous driving systems: Open challenges and emerging paradigms,” 2026.
- [135] W. Zuo, Y. Liu, W. Yang, F. Wu, W. Zhou, J. Liu, P. Sun, J. Cheng, D. Yang, and X. Zeng, “Embodied multimedia: When multimedia meets embodied intelligence,” 2026, SSRN preprint.
- [136] C. Sima, K. Renz, K. Chitta, L. Chen, H. Zhang, C. Xie, J. Beißwenger, P. Luo, A. Geiger, and H. Li, “DriveLM: Driving with graph visual question answering,” in *ECCV*. Cham: Springer Nature Switzerland, 2024, pp. 256–274.
- [137] X. Tian, J. Gu, B. Li, Y. Liu, Y. Wang, Z. Zhao, K. Zhan, P. Jia, X. Lang, and H. Zhao, “DriveVLM: The convergence of autonomous driving and large vision-language models,” in *CoRL*, ser. Proceedings of Machine Learning Research, vol. 270. PMLR, 2025, pp. 4698–4726.
- [138] L. Wen, D. Fu, X. Li, X. Cai, T. Ma, P. Cai, M. Dou, B. Shi, L. He, and Y. Qiao, “DiLu: A knowledge-driven approach to autonomous driving with large language models,” in *ICLR*, 2024.
- [139] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, “Learning transferable visual models from natural language supervision,” in *ICML*, ser. Proceedings of Machine Learning Research, vol. 139. PMLR, 2021, pp. 8748–8763.
- [140] Q. Song, H. Ye, M. Harman, and F. Sarro, “Generative AI for testing of autonomous driving systems: A survey,” *ACM Trans. Softw. Eng. Methodol.*, 2026.
- [141] W.-J. Chang, F. Pittaluga, M. Tomizuka, W. Zhan, and M. Chandraker, “SAFE-SIM: Safety-critical closed-loop traffic simulation with diffusion-controllable adversaries,” in *ECCV*, 2023.
- [142] Y. Mei, T. Nie, J. Sun, and Y. Tian, “Seeking to collide: Online safety-critical scenario generation for autonomous driving with retrieval augmented large language models,” in *ITSC*, 2025.
- [143] L. Baraldi, Z. Zeng, C. Zhang, A. Nayak, H. Zhu, F. Liu, Q. Zhang, P. Wang, S. Liu, Z. Hu, A. Cangelosi, and L. Baraldi, “The safety challenge of world models for embodied AI agents: A review,” 2025.



Han Huang is currently pursuing the M.S. degree in electronic information with the College of Electronic and Information Engineering, Tongji University, Shanghai, China. He received the B.S. degree in transportation engineering from Central South University, Changsha, China, in 2025. His current research interests include autonomous driving, world models, and intelligent transportation systems.



Dingkan Yang received the B.E. degree (with honors) in Communication Engineering from the joint training program of Yunnan University and the Chinese People's Armed Police (PAP), Kunming, China, in 2020, and the Ph.D. degree in Computer Science from Fudan University, Shanghai, China, in 2025. His research interests include multimodal learning, generative AI, and embodied AI.

Dr. Yang has published multiple papers as the first author in reputable journals and top international conferences, such as IEEE TPAMI, TCSVT, NeurIPS, CVPR, ICCV, ECCV, and AAAI. He previously worked as a Principal Researcher at ByteDance, where his team developed a series of renowned multimodal foundational models, including SAIL-VL and SAIL-Embedding. He currently serves as CTO of Fysics AI, an artificial intelligence company dedicated to developing omnimodal interaction and physical real-world perception. He is also a postdoctoral fellow at MMLab of the Chinese University of Hong Kong.



Lulu Guo (Member, IEEE) received the B.S. degree in vehicle engineering and the Ph.D. degree in control engineering from Jilin University, Changchun, China, in 2014 and 2019, respectively.

He is currently a Research Professor at Tongji University, Shanghai, China. Before joining Tongji University, he was a Postdoctoral Research Associate with the University of Georgia, Athens, GA, USA. His current research interests include advanced vehicle control, energy management, and vehicle cybersecurity.



Jing Cheng is currently an Associate Researcher with the School of Physics, East China Normal University, Shanghai, China. She received the Ph.D. degree in Electronics and Information at Fudan University, Shanghai, China, in 2022. Her research interests include fiber optic sensing and artificial intelligence.



Yang Liu (Member, IEEE) is an Assistant Professor with Tongji University, Shanghai, China, and the Leader of Embodied Perception and Intelligence Computing (EPIC) Lab, Tongji University. He received the B.S. degree (with honors) in automation from Nanjing University, Nanjing, China, in 2020, and the Ph.D. degree in computer science from Fudan University, Shanghai, China, in 2025. From 2023 to 2025, he was a Visiting Scholar with the University of Toronto, Toronto, ON, Canada, under a joint Ph.D. program. His research focuses on embodied perception with applications in autonomous driving systems.

Dr. Liu has published multiple first-author papers in prestigious journals, including ACM CSUR, IEEE TIP, IEEE TII, and IEEE TIE. He has served as the Associate Editor for Discover Computing, Guest Editor for IEEE TCSS, Area Chair for IEEE ICME 2026, BMVC 2026, and IEEE ICIP 2025, Workshop/Special Session Co-Chair for IEEE ICASSP, IEEE ICIP, IEEE MMSP, and a TPC member for top-tier conferences, including CVPR, ICCV, ICLR, NeurIPS, AAAI, and ACM MM. He is an active reviewer for ACM CSUR, PR, INFU, and multiple IEEE Transactions, including TPAMI, TIP, TII, TNNLS, TCSVT, TMM, TIFS, and TITS.



Victor C.M. Leung (Life Fellow, IEEE) is the Dean of the Artificial Intelligence Research Institute at Shenzhen MSU-BIT University (SMBU), China, a Distinguished Professor of Computer Science and Software Engineering at Shenzhen University, China, and also an Emeritus Professor of Electrical and Computer Engineering and Director of the Laboratory for Wireless Networks and Mobile Systems at the University of British Columbia (UBC), Canada. His research is in the broad areas of wireless networks and mobile systems. His published works

have together attracted more than 70,000 citations. He has been named in the Clarivate Analytics list of "Highly Cited Researchers" for several years.

Prof. Leung is serving on the editorial boards of the IEEE Transactions on Green Communications and Networking, IEEE Transactions on Computational Social Systems, and several other journals. He received the 1977 APEBC Gold Medal, 1977-1981 NSERC Postgraduate Scholarships, IEEE Vancouver Section Centennial Award, 2011 UBC Killam Research Prize, 2017 Canadian Award for Telecommunications Research, 2018 IEEE TCGCC Distinguished Technical Achievement Recognition Award, and 2018 ACM MSWiM Reginald Fessenden Award. He co-authored papers that were selected for the 2017 IEEE ComSoc Fred W. Ellersick Prize, 2017 IEEE Systems Journal Best Paper Award, 2018 IEEE CSIM Best Journal Paper Award, and 2019 IEEE TCGCC Best Journal Paper Award.

Prof. Leung is a Life Fellow of IEEE, and a Fellow of the Royal Society of Canada (Academy of Science), a Fellow of Canadian Academy of Engineering, and a Fellow of Engineering Institute of Canada.



Hong Chen (Fellow, IEEE) received the B.S. and M.S. degrees in process control from Zhejiang University, China, in 1983 and 1986, respectively, and the Ph.D. degree in system dynamics and control engineering from the University of Stuttgart, Germany, in 1997.

She joined Jilin University of Technology, China, in 1986. From 1993 to 1997, she was a Wissenschaftlicher Mitarbeiter with the Institut fuer Systemdynamik und Regelungstechnik, University of Stuttgart. Since 1999, she has been a professor at Jilin University and hereafter a Tang Aoqing professor. From 2015 to 2019, she served as the Director of the State Key Laboratory of Automotive Simulation and Control. In 2019, she joined Tongji University as a distinguished professor. Her current research interests include model predictive control, nonlinear control and applications in mechatronic systems focusing on automotive systems.

Prof. Chen is a Fellow of the Chinese Association of Automation and a Fellow of the Institute of Electrical and Electronics Engineers.